

UNIVERSITY OF HOUSTON

FOUNDATIONS OF SECURITY

COSC 6347

---

## Final Exam Review

---

*Author*

K.M. HOURANI

*Based on Notes By*

Dr. Aron LASZKA

December 7, 2020

# 1 Advanced crypto

## 1.1 Advanced Cryptographic Primitives

- secure multiparty computation
- homomorphic encryption

## 1.2 Commitment Schemes

### Commitment Problem

- Bob “calls” the coin flip (i.e., heads or tails)
- Alice flips the coin
- Bob wins if her call is correct, Alice wins otherwise

Can we prevent Alice from cheating even if the players are not in the same physical location?

### Commitment Scheme

- Two phases
  1. commit:  $A$  chooses a value  $V$ ,  $A$  sends a **commitment** of  $V$  to  $B$
  2. reveal:  $A$  reveals the value of  $V$
- Example: coin flipping
  1. commit:  $A$  flips a coin,  $A$  sends a commitment (i.e., coin is heads or tails) to  $B$ 
    - $B$  calls the coin flip (i.e., heads or tails)
  2. reveal:  $A$  reveals the value of the coin flip
- Requirements for commitment scheme
  - $B$  cannot learn the value of  $V$  from the commitment
  - $A$  can reveal only the originally chosen value for a commitment

### Naive Attempt Using Hash Function

- $H$ : cryptographic hash function
- If the set of possible values of  $V$  are small (e.g., “heads” or “tails”),  $B$  can learn  $V$  by simply trying all possible values

### Secure Commitment Using Hash Function

- Collision-free hash function  $\rightarrow A$  cannot cheat by finding  $V_1$  and  $V_2$  such that  $H(r_1 \mid r_2 \mid V_1) = H(r_1 \mid r_2 \mid V_2)$ 
  - $r_1$  prevents pre-computation of colliding  $V_1$  and  $V_2$

## 1.3 Secret Sharing

- Problem: distribute a secret among  $N$  participants such that
  - any group of at least  $T$  participants can reconstruct the secret
  - no group of fewer than  $T$  participants can reconstruct any part of it
- Types
  - **unconditionally secure**: information-theoretically secure (unbounded attacker)
  - **conditionally secure**: typically more efficient

### Special Case: $T = N$

- Unconditionally secure scheme:
  1. let the secret be a binary number  $S$
  2. pick  $N - 1$  random numbers  $R_1, R_2, \dots, R_{N-1}$  of the same length
  3. give each participant  $i$ ,  $i < N$ , the number  $R_i$


- 4. give the last participant the result of  $S \oplus R_1 \oplus R_2 \oplus \dots \oplus R_{N-1}$
- $N$  participants can reconstruct the secret by XORing their numbers:  
—

$$\begin{aligned}
 & R_1 \oplus R_2 \oplus \dots \oplus R_{N-1} \oplus (S \oplus R_1 \oplus R_2 \oplus \dots \oplus R_{N-1}) \\
 &= (R_1 \oplus R_1) \oplus (R_2 \oplus R_2) \oplus \dots \oplus (R_{N-1} \oplus R_{N-1}) \oplus S \\
 &= S
 \end{aligned}$$

- $N-1$  participants can compute only  $S \oplus R_i$ , where  $i$  is the missing participant (or  $R_1 \oplus R_2 \oplus \dots \oplus R_{N-1}$  if the last participant is missing)

## 1.4 Shamir's Secret Sharing

- Proposed by Shamir in 1979
- Unconditionally secure
- Special case  $T = 2$ 
  1. let the secret be a number  $S$
  2. pick random number  $A$
  3. let each participant's share be a random point on the line  $Ax + S$



- $T = 2$  participants can reconstruct the secret since any two points define a line
- Single participant cannot learn the slope

### General Case

- Arbitrary  $T$ :
  1. let the secret be a number  $S$
  2. pick random numbers  $A_1, A_2, \dots, A_{T-1}$
  3. let each participant's share be a random point from the curve

$$y = S + A_1x + A_2x^2 + \dots + A_{T-1}x^{T-1}$$

- At least  $T$  points are necessary to define a polynomial of degree  $T - 1$
- Example  $T = 3$ 
  - secret is a parabola (i.e.,  $A_2x^2 + A_1x + S$ )
  - there an infinite number of parabolas fitting two points
  - but three point define one uniquely

## 1.5 Secure Multiparty Computation

- Problem:  $N$  participants with private data  $d_1, d_2, \dots, d_N$ 
  - participants would like to compute the value  $F(d_1, d_2, \dots, d_N)$  of a public function  $F$  over their private data

- no participant  $i$  would like to reveal any information about its data  $d_i$
- Requirements
  - **privacy**: no information is revealed about any private data (other than what is revealed by the public output)
  - **correctness**: public function is correctly computed
- Adversaries may be semi-honest (passive) or malicious (active)

## 2 WiFi security

### 2.1 Security Challenge

- **Problem**: no inherent physical protection
- **joining** a network does not require physical access
- radio transmissions are broadcast → anyone in range can **eavesdrop**
- **injecting** new messages or **replaying** old messages is possible
- **jamming** attacks against availability
- jamming and injecting messages can be combined into **tampering** attacks

### 2.2 Simple “Solutions” for Access Control

#### Hidden SSID

- Association request must contain the SSID of the network
  - by default, the AP broadcasts it periodically in the beacon
- AP may be configured to **stop announcing the SSID** → SSID may be used as a “password”
- However,
  - SSID must be hard to guess
  - every authorized user must know the SSID
  - **SSID can be easily eavesdropped** whenever an authorized station connects to the network → does not provide any security
- Tools are available for eavesdropping (e.g. Aircrack-ng)

#### MAC Address Based Filtering

- AP may be configured to **allow only devices with certain MAC addresses** to connect
  - MAC addresses of all authorized devices must be registered in advance
- However,
  - **MAC address is sent in plaintext** in every packet
  - many WLAN devices allow their MAC addresses to be changed → attacker can easily impersonate an authorized user

### 2.3 802.11 Security Standards

#### WEP

- security is based on a 40 or 104-bit secret key
  - WiFi “password” shared by all users
- confidentiality: RC4 stream cipher
  - key is extended by a 24-bit IV, which is changed for each message → used as nonce to prevent key reuse problems
- integrity: encrypted CRC32 (Cyclic Redundancy Check) checksum
- access control: challenge-response between AP and station

## WEP Design Flaws

- Authentication
  - **one-way authentication** (only for station) → AP can be impersonated
- Integrity protection
  - based on **error-detection code** (CRC32) instead of cryptographic hash → forging authentication tags is trivial
  - **no message replay protection**
- Key usage
  - **no session key**: long-term key used for all purposes (authentication, encryption, integrity protection)
  - **short nonce** (i.e., 24-bit IV) → danger of key reuse for stream cipher
    - \* busy network with 1000 packets per second reuses in less than 5 hours
  - vulnerable to Fluhrer-Mantin-Shamir Attack
    - \* In practice, WEP keys can be broken in a matter of minutes (or less) → WEP is **not secure**

## 2.4 WiFi Protected Access (WPA)

### WPA

- Standard: 802.11i TKIP (Temporal Key Integrity Protocol)
- Design goals: **fix the flaws of WEP** and be **compatible with legacy hardware**
- Overview
  - key usage: session key is established during a secure two-way authentication
  - confidentiality: RC4 encryption, but with **48-bit IV**, which is **mixed thoroughly** with the session key and source MAC address
    - \* prevents key reuse and the Fluhrer-Mantin-Shamir attack
  - integrity: 64-bit message integrity codes computed using Michael, which is **computationally very efficient** but provides only 20 bits of effective security
    - \* after wrong code, station is banned for a minute and needs to re-authenticate
  - Deprecated in later revisions of the standard
- 

### WPA-2

- Standard: **IEEE 802.11i**
- WPA 2 Devices can be certified by the Wi-Fi Alliance

### Phases

1. Discovery
  - agree on what authentication method and ciphers to use
2. Authentication
  - may use an authentication server
  - create a master session key
3. Key management
  - derive keys for various purposes
4. Protected data transfer
5. Connection termination

## Discovery Phase

- Goal: station and AP may support different sets of authentication methods and ciphers → they need to agree on which ones they will use
- **Authentication and key-management suite**: how to perform mutual authentication and derive fresh keys
  - IEEE 802.1X, pre-shared key (PSK), or vendor-specific
- **Cipher suite**: what ciphers to use for confidentiality and integrity
  - WEP, TKIP, CCMP, or vendor-specific
- Protocol
  1. AP can periodically **broadcast** its security capabilities using a **Beacon** (or station can ask for it using a Probe Request message)
  2. Station **specifies** an authentication and cipher suite in an **Association Request**
  3. if the AP **accepts** the specified suites, it sends an **Association Response**

## Authentication Phase

- Goals:
  - mutual authentication:
    - 1) only authorized stations can use the network
    - 2) station is assured that it communicates with a legitimate network
  - generate **pairwise master key** (PMK)
- Approaches
  - Pre-shared key (PSK)
    - \* password is deployed on each station and the AP manually
    - \* PMK = PSK = generated from the password using a hash function
    - \* ideal for home and small office networks
  - IEEE 802.1X

## Key-Management Phase

- Goals:
  - derive **pairwise transient keys** from the PMK
  - distribute **group keys**
- Pairwise transient key (PTK)
  - protecting data between station and AP
  - generated from PMK and the AP's and station's MAC addresses and nonces
- Group temporal key (GTK)
  - protecting multicast communication
  - group master key (GMK): generated randomly by the AP
  - distributed using the PTK

## Protected Data Transfer Phase

- Standard defines two schemes: TKIP and CCMP
- TKIP: see WPA
- CCMP (Counter mode CBC-MAC Protocol)
  - based on the CCM (Counter with CBC-MAC) authenticated encryption mode
  - integrity: CBC-MAC based on AES encryption
  - confidentiality: AES encryption in counter (CTR) mode
  - same 128-bit key for integrity and confidentiality (from PTK)
  - 48-bit packet number to prevent replay attacks

## IEEE 802.1X

- Standard for port-based network access control
- Entities
  - supplicant = station
  - authenticator = access point
  - authentication server
- Port-based: supplicant can access only the authentication server until the authentication succeeds
- Authentication server does not have to be implemented on the access point → little overhead for the access point

## EAP Authentication Methods

- Extensible framework, not a specific authentication mechanism
- Example methods
  - EAP-TLS: based on public-key certificates
  - EAP-GPSK (Generalized Pre-Shared Key): based on secret keys shared by the client and the server, uses symmetric-key cryptography

## 3 IPSec

- Collection of protocols and mechanisms, standardized by the Internet Engineering Task Force (IETF) in a series of publications
- Provides
  - data confidentiality and integrity
  - source authentication (prevent address spoofing, i.e., sending from fake address)
  - protection against packet replay
- Below the transport layer (TCP or UDP) → transparent to applications
- End-to-end security between two hosts, a host and a network, or between two networks
- Example Applications of IPSec
  - Secure remote access over the Internet
  - Secure virtual private network

### 3.1 Transport Mode and Tunnel Mode

- Transport mode
  - protects the payload of the IP packet
  - typically host-to-host communication
- Tunnel mode
  - protects the entire IP packet by encapsulating it in the payload of a new IP packet
  - typically host-to-network or network-to-network communication

Protocol		
	Authentication Header (AH)	Encapsulating Security Payloads (ESP)
Modes	both transport and tunnel	
Provides	integrity, replay protection	integrity, confidentiality, replay protection
Protects	payload and IP header	payload

## Authentication Header

- Services
  - data and origin integrity
  - replay-prevention
- Message authentication

- computed from immutable fields of the IP header, AH header (except ICV), and original payload
- algorithms: HMAC-MD5, HMAC-SHA-1, HMAC-SHA-2, ...

### Encapsulating Security Payload

- Services: confidentiality, integrity (optional), replay prevention
- Encryption: AES-CBC, 3DES-CBC, ...
- Message authentication: HMAC-SHA-1, AES-GMAC, ...
- Authenticated encryption: AES-GCM

### Combining Modes and Protocols

- Tunnel mode advantage: requires support only at the gateways
- Transport mode advantage: requires support only at the hosts
- AH advantage: authenticates some elements of the original header
- ESP advantage: protects both integrity and confidentiality
- Combining modes
  - IPSec tunnel can carry any IP packet → IPSec transport or tunnel packets can be sent through an IPSec tunnel
  - IPSec transport can protect any IP packet → IPSec transport or tunnel packets can be protected by outer IPSec transport
  - ...
  - can be nested to any depth
- Combination Examples
  1. AH in transport (for integrity) + ESP in transport (for confidentiality)
  2. IPSec packets over tunnel

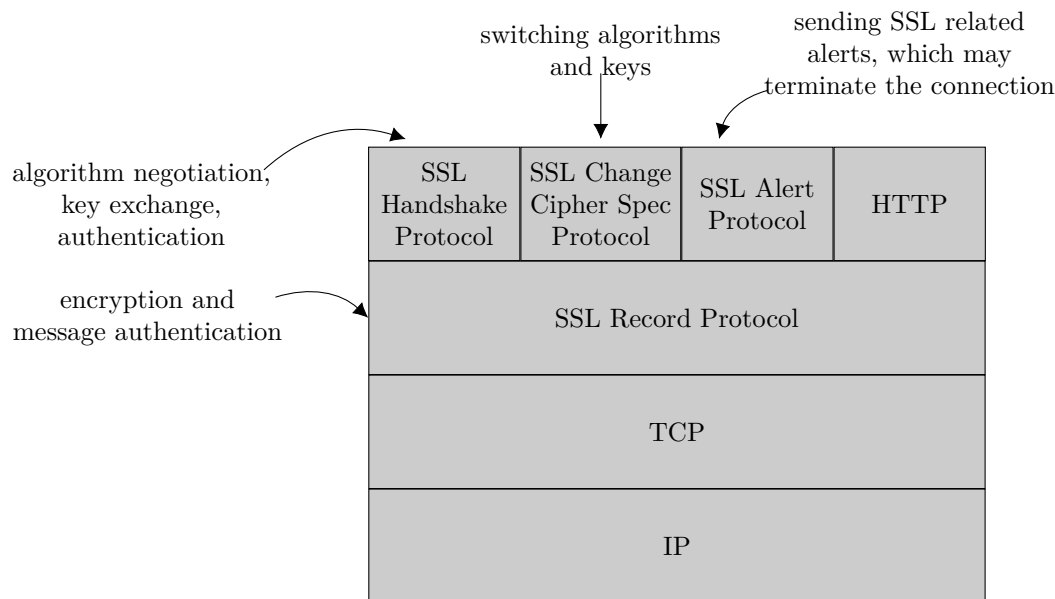
## 4 SSL / TLS

### Secure Socket Layer

- End-to-end security between two applications
- Endpoint applications can implement it without the help of the operating systems or any intermediate devices
- Developed by Netscape for securing HTTP → HTTPS = HTTP over SSL
- very widely used, not just for HTTP (e.g., FTP, POP3, IMAP)



## Overview



### SSL Record Protocol

- Security 528
  - confidentiality: symmetric-key encryption (AES GCM, Salsa20, ...)
  - integrity: message authentication codes based on symmetric-key cryptography (HMAC-SHA256, ...)
- Additional Services
  - fragmentation: fragment application data into records of at most 16 KiB
  - lossless compression: optional (default is no compression)

### SSL Handshake Protocol

- Phase 1: establish security capabilities
  - client\_hello:
    - \* highest SSL version supported by the client
    - \* nonce (timestamp + random value)
    - \* cipher suite: list of key-exchange methods, as well as encryption and MAC algorithms
    - \* compression method: list of supported compression algorithms
  - server\_hello:
    - \* highest SSL version supported by both the client and the server
    - \* nonce
    - \* chosen cipher suite and compression method
- Phase 2: server authentication and key exchange
  - certificate (optional): X.509 certificate (may be a chain)
  - server\_key\_exchange (optional):
    - \* parameters for Anonymous or Ephemeral Diffie-Hellman exchange
    - \* public-key for RSA exchange if the certificate contains only a signing key
    - \* signed by the server (together with the nonces)
  - certificate\_request (optional): ask client for an X.509 certificate
  - server\_hello\_done: server is finished
- Phase 3: client authentication and key exchange
  - certificate (optional): X.509 certificate if the server asked for a client certificate
  - client\_key\_exchange:
    - \* pre-master secret encrypted using public RSA key of the server

- \* parameters for D-H exchange
- certificate\_verify (optional): digital signature of all previous handshake messages

### Key Exchange Methods

- Goal: exchange or agree on a pre-master secret (PMS)
- RSA
  - client generates pre-master secret (PMS)
  - sends PMS to the server encrypted using RSA public-key encryption
- Diffie-Hellman protocol
  - anonymous D-H: basic D-H with no authentication
  - fixed D-H: D-H parameters of the server ( $X_A$  and  $Y_A$ ) are fixed and  $Y_A$  is contained in a digital certificate
  - ephemeral D-H: D-H with authentication

### SSL Change Cipher Spec Protocol

- Same for client and server
  - change\_cipher\_spec: signals that the communication party is switching to the negotiated cryptographic algorithms and keys
  - finished: hash value computed from the master secret and all handshake messages using HMAC with SHA hash function

### Session Resume

- Authentication and key exchange are complex → result needs to be reusable
- Session
  - association between a client and a server
  - cipher suite, compression method, and master secret
- Connection
  - within a session
  - keys and IVs for encryption and message authentication
- Session ID: identifies a session
  - sent in ClientHello → may specify an existing session to be resumed
  - sent in ServerHello → server can accept resume by sending the same ID
- Session resume skips all messages in the Handshake after the ClientHello and ServerHello messages
  - new keys and IVs are generated from the nonces in the Hello messages

### Transport Layer Security

- only minor differences compared to SSL 3.0:
  - pseudorandom function for generating keys and MAC is based on HMAC
  - variable length padding (may prevent traffic analysis)
  - other minor changes

### HTTPS

- HTTP over SSL/TLS
- Between the web browser and server:
  - HTTP client = SSL client
  - HTTP server = SSL server
- Conventions

- URL: **https://** instead of **http://**
- default TCP port is 443 instead of 80
- HTTP request may be sent after SSL Finished messages
- Protected information
  - URL, contents of the document, browser forms, cookies, headers
  - page served over HTTPS may include elements retrieved using HTTP

## 5 DNSSEC

### Domain Name System (DNS)

- Millions of domain names → distributed database
  - dynamic data ← IP address for a host may change
  - decentralized authority ← each name has an owner
- Hierarchical name space
  - domain: node in the DNS tree
  - for each domain, there is an authoritative server
- Authoritative server
  - responds to queries about the domains for which it is responsible
  - may refer to other authoritative servers for a subdomain

### DNS Queries and Responses

- Transport protocol
  - UDP port 53
  - TCP for long responses (and some tasks between nameservers)
- Messages: query and reply
  - 16-bit identification field: match queries with replies
- Caching
  - received responses are cached by servers
  - each record has a Time-to-Live field, after expiry it must be queried again

### DNS Weaknesses

- DNS responses are not authenticated
  - responses can be sent over UDP transport protocol → anyone can respond from a spoofed (i.e., fake) IP address to a query
- DNS is a key infrastructure
  - resolvers trust responses, users trust resolvers
  - by tampering with responses, an attacker may direct users to malicious websites or direct e-mail to malicious servers
- DNS cache poisoning
  - attacker sends malicious response to a DNS server, which caches it → malicious response is served to all clients using the server
  - attacker does not have to be man-in-the-middle
- Race to Respond First
  - If the attacker responds before the authoritative server, the DNS cache is poisoned with the malicious IP address
    - \* fake response is a single UDP packet from the spoofed (i.e., forged) IP address of the authoritative server


### Domain Name System Security Extensions (DNSSEC)

- Set of extensions defined by the IETF to the DNS protocol

- Guarantees origin authenticity and data integrity of DNS replies
- Backwards compatible: responses can be interpreted by DNS servers and clients that do not support DNSSEC
  - of course, no guarantees are provided for these servers and clients
- Does not provide confidentiality
  - responses are only authenticated but not encrypted
- Based on public-key cryptography
  - every response is digitally signed

### DNSSEC Public Keys and Signatures

- Signature algorithms: RSA-SHA1, RSA-SHA256, ECDSA-SHA256, ...
- Trust anchor
  - known public key for an authoritative nameserver
  - typically included in the operating systems
- Authentication chain



- Responses may include
  - RRSIG: digital signature for the contents of the response
  - DNSKEY: public key for a zone ← if the response delegates

### DNS over HTTPS and TLS

- DNS over HTTPS (DoH)
  - DNS queries and responses are encoded into HTTPS requests and responses
    - \* provides integrity and confidentiality
    - \* server may use HTTP/2 Push to send records in advance
  - controversies and criticism: can impede traffic analysis for cybersecurity, may provide false sense of privacy, can impede traffic filtering by ISPs, ...
- DNS over TLS (DoT)
  - provides integrity and confidentiality

## 6 SSH

### Secure Shell (SSH) Protocol Stack

<b>SSH User Authentication Protocol</b> Authenticates the client-side user to the server.	<b>SSH Connection Protocol</b> Multiplexes the encrypted tunnel into several logical channels.
<b>SSH Transport Layer Protocol</b> Provides server authentication, confidentiality, and integrity. It may optionally also provide compression.	
<b>TCP</b> Transmission control protocol provides reliable, connection-oriented end-to-end delivery.	
<b>IP</b> Internet protocol provides datagram delivery across multiple networks.	

#### SSH Transport Layer

- Packet Exchange
- Identification string exchange
  - both the client and the server send: `SSH-protocolVersion-softwareVersion`
- Algorithm negotiation (KEXINIT)
  - both the client and the server send the list of key exchange, encryption, MAC, and compression algorithms that they support
  - chosen one: first on the client's list that is also on the server's list
- End of key exchange (NEWKEYS)
  - start using the algorithms and keys

#### Server authentication

- servers signs a hash of all the earlier messages and the new symmetric key
- servers sends the signature and its public key to the client
- Client needs to verify the public-key sent by the server (i.e., verify that the public key belongs to the server host)
- Trust models
  - Certificate authority
    - \* client accepts public keys that are certified by a trusted CA
  - Local database
    - \* client has a list of known pairs of hosts and public-keys
    - \* typically, each user has a list stored in its home directory
      - default location: `~/.ssh/known_hosts`
- Known Hosts
  - First connection: verify and store host and public key
  - Subsequence connections: compare to stored key

## SSH User Authentication Methods

- Passwords
  - client sends a username and a password
- Public key
  - client sends a public key and a signature based on the corresponding private key
  - server checks if the public key is acceptable and verifies the signature
  - typically, for every user account on the server host, there is a list of acceptable public keys stored at `~/.ssh/authorized_keys`
- Host based
  - assumes that the server trusts the client host → since the client host has already authenticated the user, the server only needs to verify the identity of the client host
  - client sends a signature based on the private key of the client host

## 7 E-mail security

### Weaknesses

- Simple Mail Transfer Protocol
  - Outgoing emails can be tampered (integrity), eavesdropped (confidentiality), or forged (authenticity/integrity)

### E-Mail Threats

- Spam: unsolicited messages sent in bulk
- E-mail scam: advance-fee scam (a.k.a. “Nigerian Prince” scam), job scam, ...
- Phishing: collecting sensitive information (e.g., passwords, credit card numbers) or delivering malware by impersonating a trusted entity
- Spear-phishing: phishing directed at specific targets (e.g., users)


### Pretty Good Privacy (PGP)

- General-purpose application for secure communication between users
  - confidentiality and integrity protection for files and e-mail
  - built on widely used asymmetric and symmetric-key cryptographic algorithms
  - communicating users know each other’s public keys → trust
- IETF standard: OpenPGP
- Software
  - PGP went commercial in 1996
  - GnuPG is a free and open-source implementation of OpenPGP

### PGP Key Management

- Each user may have multiple public-private key pairs → key identifiers are used to specify which key is used
- Key storage
  - private-key ring: user’s own public-private key pairs
    - \* each entry has user identifier, key identifier, public key, encrypted private key
    - \* private key is encrypted using a passphrase
  - public-key ring: public keys of other users
    - \* each entry has user identifier, key identifier, public key, trust levels, and signatures from other users
    - \* public-keys can be verified directly (e.g., delivery on secure channel) or using the “web of trust”

## PGP Web of Trust



## PGP Authentication

- Digital signature using the sender's private key
  - hash using MD5, SHA-1 or SHA-2, and then sign using RSA or DSA
- Message may be compressed after signature
  - ZIP, Bzip2, ...

## PGP Encryption

- Message may be compressed before encryption
  - ZIP, Bzip2, ...
- Generate a new 128-bit random symmetric key for each message
  - encrypt the message with the symmetric key using a block cipher in CFB mode (3DES, Blowfish, AES, ...)
  - encrypt the symmetric key with the recipient's public key using RSA or ElGamal

## S/MIME

- MIME (Multipurpose Internet Mail Extension)
  - fixes the limitations posed by SMTP
  - New headers fields
    - \* Content-Type: type of message content
      - multipart type: body contains multiple parts, each having a header (e.g., images in HTML message, attachments)
      - simple types (e.g., text/plain, image/jpeg, text/html)
    - \* Content-Transfer-Encoding: how binary data is represented in 7-bit ASCII (e.g., Base64, quoted-printable)
- Secure / Multipurpose Internet Mail Extension (S/MIME)
  - security enhancement to the MIME e-mail format standard
  - similar to PGP (Pretty Good Privacy)
    - \* both S/MIME and PGP enable encrypting and signing messages
    - \* both have IETF standards
    - \* both support state-of-the-art algorithms (AES, RSA, SHA-2, ...)
    - \* S/MIME is likely to emerge as the industry standard for commercial and organizational use (e.g., Microsoft Outlook and Gmail support S/MIME)
    - \* PGP is likely to remain the choice for personal e-mail security

## S/MIME Functionality

- Functions
  - Signed data: message is digitally signed, and both the signature and the message are encoded (using Base64 representation)
  - Clear-signed data: similar to signed data, but only the signature is encoded

- Enveloped data: encrypted message content and encrypted content-encryption key (i.e., session key) for one or more recipients (encoded using Base64)
- Signed and enveloped data: signing and encrypting may be nested
- MIME content types
  - application/pkcs7-mime: signed or enveloped data
  - multipart/signed: for clear-signed data, which contains a message and a signature (signature part is application/pkcs7-signature)

## **S/MIME Public Keys**

- Public-key certificates
  - based on X.509 digital certificate format
  - similar to PGP, digital certificates are distributed manually
  - however, certificates may be signed by a CA
- Public keys are used for
  - verifying signatures
  - encrypting session keys
    - \* for enveloped data, the sender generates a random session key
    - \* session key is encrypted with each recipient's public key, and the message contents are encrypted (using symmetric-key crypto) with the session key
    - \* upon receiving the message, a recipient can decrypt the session key, and then decrypt the message contents using the session key

## **DKIM**

- E-Mail Spoofing Problem
- Attackers may use e-mail with forged sender addresses for spam/phishing
- Limitations of PGP and S/MIME
  - depend on the sending and receiving users, who must install or configure software, share public keys, etc.
  - do not sign the message header, only the message contents
- DomainKeys Identified Mail (DKIM)
  - specification for signing e-mail messages
  - implemented on the servers, therefore transparent to the users

## **DKIM Signature and Verification**

- Signature: Dkim-Signature header field
  - v: version
  - a: algorithm used for signature (RSA-SHA1, RSA-SHA256)
  - d: domain name
  - s: selector (if there are multiple public keys)
  - h: list of signed header fields
  - bh: hash of the body part of the message
  - b: signature in Base64
- Verification
  - receiving SMTP server queries record s.\_domainkey.d using DNS
  - nameserver returns the public key corresponding to the signing private key

## **E-Mail Authentication Solutions**

- DomainKeys Identified Mail (DKIM)



- Sender Policy Framework (SPF)
  - another system for preventing e-mail spoofing
  - DNS record lists all authorized sending hosts (i.e., e-mail servers) for a domain → receiving server can verify
  - may be combined with DKIM
- Domain-based Message Authentication, Reporting and Conformance (DMARC)
  - built on top of DKIM and SPF, published using DNS record
  - enables domain owners to publish a policy (combination of DKIM and SPF) for verifying the legitimacy of e-mails from the domain

## 8 Authentication and access control

### Authentication

- Authentication: reliably verifying the identity of someone or something
  - computer authenticates another computer
  - computer authenticates a user
- Types of authentication
  - one-way authentication or mutual authentication
  - one-time or establishing a session (e.g., combined with key exchange)
- Typical methods of computer authentication
  - cryptography-based
  - address-based

### User Authentication

- Types of user authentication factors
  - knowledge: some secret known only by the user (e.g., password)
  - ownership: some physical object possessed by the user (e.g., bank card)
  - inherence: some physical characteristic of the user (e.g., fingerprint)
- Password-based user authentication
  - typical form of knowledge-based authentication
  - verifier stores the password in a database or file
  - often combined with cryptography-based approaches to protect the password from eavesdropping
  - password must be easy to remember but hard to guess

### Password-Based Authentication

- Problem: easy-to-remember passwords are weak
  - Miller's law: number of objects an average human can hold in working memory is  $7 \pm 2$
  - length of passwords that users can easily remember (i.e., not write down somewhere) is very limited
- Brute-force attack: password guessing
  - online: attacker must rely on the verifier to test the correctness of a password → verifier can limit the number of attempts (e.g., number of unsuccessful login)
  - offline: attacker can test the correctness of a password on its own

### Password Storage

- Cleartext passwords are insecure
  - system administrators (and other local users) may easily read passwords
  - attackers who have compromised a system may be able to read passwords

- Users tend to reuse passwords → breach may affect other systems as well
- Store the cryptographic hash of the password
  - during authentication, the user enters the plaintext password, and the verifier computes its hash and compares it with the stored hash
  - attacker can perform offline guessing to recover the plaintext password
- Brute-forcing multiple hashed passwords
  - first, precompute a table of [password, hash] values for possible passwords
  - second, for each hashed password, look up the precomputed hash value

### Salting

- Before hashing a password mix it with a salt value
  - both when the password is set and during verification
  - verifier stores: username, salt,  $H(\text{password} + \text{salt})$
- Salt value
  - randomly generated for each user account
  - may be stored in plaintext by the verifier
- Salt values do not have to be memorized → strong randomness
  - prevents precomputing hashes since the attacker cannot consider all possible salt values (different salt values require different precomputation)
  - also hides identical passwords, which would result in identical hashes
- However, it does not make guessing a single password harder (assuming that the attacker knows the salt)

### Multi-Factor Authentication

- User is authenticated only after passing multiple independent authentication mechanisms
  - typically, each mechanism is built on a different type of factor (e.g., knowledge + possession), so it is independent of the other mechanisms
  - attacker must circumvent all authentication mechanisms to succeed
- Possession factors
  - disconnected token: not connected to the client computer, typically the user manually enters authentication data displayed by the token
  - connected token: physically connected to the client computer (e.g., USB token)
- Inherence factors
  - includes fingerprint, face, voice, or iris recognition

### Access Control

- Access control (i.e., authorization): approving or rejecting access requests
- Abstractions
  - subjects: entities that can perform actions on the system
  - objects: resources to which access must be controlled
- Control access to objects based on a policy

### Discretionary Access Control (DAC)

- Allows access rights to be propagated at the subjects' discretion
- Often implemented using the notion of owner
  - every object has an owner subject, who can set the permissions for that object
- Used by popular operating systems (e.g., Unix and Windows)
- Problem: non-malicious users are not necessarily trustworthy
  - phishing: subjects may be tricked into propagating their access rights to malicious entities

- malware: malicious code running with a subject’s credentials can disclose or modify sensitive information
- large organizations working with sensitive data may need centralized control

## **Mandatory Access Control (MAC)**

- Restricts the access of subjects to objects based on a system-wide set of rules
  - system-wide rules are set by a central authority (e.g., system administrator)
  - policy is mandatory → users do not have full control over access to the resources that they create
- Traditionally used for implementing multilevel security
  - objects have security classifications (e.g., “Top Secret”, “Secret”)
  - subjects have security clearances
- Available in some form on many modern operating systems
  - SELinux and AppArmor for Linux, and Mandatory Integrity Control for Windows
- May be combined with DAC: grant access only if both DAC and MAC permit the access

## **Access Control Models**

- Access control list (ACL): list permissions for each object
  - for each object, list pairs of [subject, access right]
- Role-based access control (RBAC): row oriented
  - create a set of roles (e.g., based on real-world job functions), and assign a role (or roles) to each subject
  - for each role, list pairs of [object, access right]

## **Unix Access Control**

- user: has a unique UID (special UID = 0 for root user)
- group (collection of multiple users): has a unique GID
- Access control abstraction
  - subject = process
    - \* has an effective UID and GID (as well as real and saved UIDs and GIDs)
  - object = file
    - \* has an owner (UID) and a group (GID), typically inherited from the process that created the file
    - \* almost everything is a file on a Unix system (regular files, directories, devices, Unix domain sockets, ...)
  - Each file has 12 permission bits
    - \* read, write, and execute permission for owner, group, and others
    - \* set user ID (setuid), set group ID (setgid), sticky bits
  - When a process wants to read/write/execute a file,
    1. if effective UID = file owner → use read/write/execute permission for owner
    2. else if effective GID = file group → use read/write/execute permission for group
    3. else → use read/write/execute permission for others
  - For directories,
    - \* read means listing the contents of the directory
    - \* write means creating, renaming, and deleting files in the directory
    - \* execute means accessing the files (and directories) within the directory (must also have execute permission on all the parent directories)

## **Sticky, Set UID, and Set GID Bits**

- Sticky bit

- when set on a directory, files within that directory can be renamed or deleted only by their owners, the directory owner, or a superuser
- for example, sticky bit is typically used on the /tmp directory
- Set UID bit
  - when set on an executable file, the effective UID of a process executing the file is set to the file owner UID
  - for example, set UID bit is typically used on the passwd command
- Set GID bit
  - when set on an executable file, the effective GID of a process executing the file is set to the file group GID
  - when set on a directory, new files created within will inherit the GID of the directory

## 9 Software vulnerabilities and countermeasures

### Buffer Overflow

- Buffer overflow / buffer overrun: anomalous condition where a process tries to store data beyond the boundaries of a fixed-length buffer
  - extra data overwrites adjacent memory, which may lead to denial-of-service or arbitrary code execution
- Vulnerable programming languages
  - C, C++ and other “lower-level” languages without bounds checking
  - “higher-level” languages, such as Java and C#, are generally not vulnerable
- Attacker can supply malicious (i.e., long) input
  - remotely through a network connection (e.g., specially crafted HTTP request)
  - locally (e.g., by sending a specially crafted file to the target user)

### Typical Buffer Overflow Exploits

- Attacker’s goal: execute arbitrary code (i.e., code chosen by the attacker)
- execute arbitrary code (i.e., code chosen by the attacker)
- Stack-based exploitation (“stack smashing”)
  - overflow a local buffer
  - overwrite local variables, function return addresses, exception handlers, etc.
- Heap-based exploitation
  - overflow a dynamically allocated buffer
  - overwrite other data, function pointers, etc.
- Shellcode: small piece of code, which allows the attacker to control the compromised machine

### Buffer-Overflow Exploit Countermeasures

- Compile-time – hardening new software
  - programming languages
  - safe functions and libraries
  - compiler extensions
- Run-time – protecting existing software
  - executable space protection
    - \* Lot of exploits build on injecting and executing malicious code
    - \* By separating the memory space of a process into executable and modifiable parts, code injection can be prevented
    - \* Problem: modern computer architectures do not separate code from data

- \* Limitation: cannot fully protect programs that create and execute code at runtime (e.g., just-in-time compilers)
- \* Circumventing Executable Space Protection
  - Attacker can re-use existing code from the memory space of the process for malicious purposes
  - Return-to-libc attack
  - for most processes, the standard C library is loaded into memory
  - attacker can change the return address of a function to point to the beginning of a function in the C library
  - common target: system function – takes as argument a string, and executes it as a system command with the privileges of the process
  - attacker has control over the stack → attacker can set up parameters for the C library function
- address space layout randomization
  - \* In order to reliably jump to an exploited code, the attacker needs to know its address
  - \* Address Space Layout Randomization (ASLR)
    - randomly arrange the positions of the executable, the stack, and the heap in the process's address space
    - may prevent return-to-libc attacks
    - most operating systems (e.g., Windows, Linux) implement some randomization
  - \* Counter-countermeasures
    - information leakage (e.g., printf vulnerability)
    - random guessing (Heap spraying)

## Integer Overflow

- Integer overflow occurs when an arithmetic operation leads to a value that is greater than the maximum value that can be stored
- Affects most languages - even some managed languages, such as Java and C#, are susceptible to integer overflow errors
- Exploiting integer overflow errors
  - calculating indexes into arrays
  - calculating the amount of space to allocate for a buffer
  - checking whether an overflow could occur

## Input Validation

- Sources of input
  - user supplied files and terminal input
  - command line arguments
  - environment variables
  - function calls from other modules
  - network packets (web applications in detail later)
  - ...

## Format String Vulnerabilities

- Vulnerable functions
  - `sprintf`: writes to buffer
  - `fprintf`: writes to file
  - other members of the printf family (e.g., `snprintf`)
  - `printf`: used in the Linux kernel
  - other functions that use format strings (e.g., `syslog`)
- Format placeholders

- `%s` – string
- `%d` – number (output in decimal format)
- `%x` – number (output in hexadecimal format)
- Special placeholder: `%n`
  - \* argument must be a pointer to a signed integer, where the number of characters printed so far will be written
- When variable is printed directly, introduces Vulnerability
  - `printf(name);`
  - if `name` is provided as `%d%d%d` then the top three stack values will be printed
- To prevent, use string formatting
  - `printf("%s%", name);`
  - if `name` is provided as `%d%d%d` then `"%d%d%d"` will be printed

## Race Conditions

- Race condition
  - when results depend on the sequence or timing of uncontrollable events
  - for example, when the output of a software depends on how the operating system schedules the execution of multiple processes or threads
- Typically happens when interacting with
  - memory shared by multiple processes (or threads)
  - file system
  - signals and other interprocess communication mechanisms
- Race condition bugs and errors
  - happen when events do not occur in the intended order
  - typically very difficult to reproduce and debug
- Attack approaches
  - try multiple times (if possible)
  - slow down the target process
  - increase computational load on the machine
  - computational complexity attacks
- Preventing Race Conditions
  - Time of check to time of use
    - \* we cannot allow any changes in this interval
    - \* trying to make it short is not enough
  - Prevention techniques
    - \* work with file descriptors instead of filenames
    - \* rely on filesystem access checks
    - \* be careful with directories that are writable by everyone (e.g., `/tmp/`)
    - \* lock resources (files, databases, etc.)
    - \* look out for non-atomic operations (e.g., `num++`)
    - \* synchronization (e.g., semaphore, mutex)

- 10 Web vulnerabilities
- 11 Malware
- 12 Secure development
- 13 Detection
- 14 Isolation
- 15 Denial of Service attacks
- 16 Vulnerability scanners